

Machine Learning 6.867 - Pset 3

November 7, 2015

1 Multi-Class SVM

2 Neural Networks

Neural networks are used in machine learning to make predictions, similar to logistic regression, SVM, or regression. We can represent neural networks using a graph with nodes and edges (see Bishop figure 5.1). Assume that we observe data $(\mathbf{x}^{(i)}, \mathbf{y}^{(i)})$, $i = 1, \dots, N$, where $\mathbf{x}^{(i)} \in \mathbb{R}^{D+1}$ and $\mathbf{y}^{(i)} \in \{0, 1\}^K$. Let $(\mathbf{x}, \mathbf{y}) = ([x_0, x_1, \dots, x_D], [y_1, \dots, y_K])$ be a general observation, where $x_0 = 1$ is a constant term for the bias. We create nodes for each of the features x_i , referred to as *inputs*, and nodes for each of the class labels y_i , referred to as *outputs*. Next, we introduce a series of nodes in the middle of the graph, called *hidden units*, and we draw edges connecting the **inputs** \rightarrow **hidden units** \rightarrow **outputs**. The key idea in neural networks is that we can model the hidden units as functions of the inputs, and model the outputs as functions of the hidden units.

For 2-layer neural networks, we have one layer of hidden units denoted by $[z_0, z_1, \dots, z_m]$. There are weights $\mathbf{w}_{ji}^{(1)}$ for each edge $x_i \rightarrow z_j$ and $\mathbf{w}_{kj}^{(2)}$ for each edge $z_j \rightarrow y_k$, which are unknown and will be learned through training the model. However, there are no edges from the inputs to z_0 , because we assume $z_0 = 1$ is a constant term for the bias. Let $\sigma(\cdot)$ denote the logistic function, which we will use as the *activation function* for both the hidden and output layers of our neural network. The predicted value for output k given parameters $\mathbf{w} = \{\mathbf{w}^{(1)}, \mathbf{w}^{(2)}\}$ and input \mathbf{x} will be:

$$h_k(\mathbf{x}, \mathbf{w}) = \sigma \left(\sum_{j=1}^M w_{kj}^{(2)} \sigma \left(\sum_{i=0}^D w_{ji}^{(1)} x_i \right) + w_{k0}^{(2)} \right). \quad (1)$$

2.1 Implementation

2.1.1 Gradient Descent

We implemented a 2-layer regularized neural network using gradient descent. The loss function, which is the negative log-likelihood, is equal to:

$$\ell(\mathbf{w}) \equiv \sum_{i=1}^N \sum_{k=1}^K \left[-y_k^{(i)} \log(h_k(\mathbf{x}^{(i)}, \mathbf{w})) - (1 - y_k^{(i)}) \log(1 - h_k(\mathbf{x}^{(i)}, \mathbf{w})) \right] \quad (2)$$

We add a regularizer term, and the final cost function becomes:

$$J(\mathbf{w}) \equiv \ell(\mathbf{w}) + \lambda(\|\mathbf{w}^{(1)}\|_F^2 + \|\mathbf{w}^{(2)}\|_F^2), \quad (3)$$

where $\|A\|_F = \sqrt{\sum_{i,j} A_{ij}^2}$ is the matrix Frobenius norm. To take the derivative of this function, we introduce some more notation:

$$a_j^{(1)}(\mathbf{x}) \equiv \sum_{i=0}^D w_{ji}^{(1)} x_i \quad (4)$$

$$a_k^{(2)}(\mathbf{x}) \equiv \sum_{j=1}^M w_{kj}^{(2)} \sigma(a_j^{(1)}(\mathbf{x})) + w_{k0}^{(2)} \quad (5)$$

We refer to $a_j^{(1)}(\mathbf{x})$ and $a_k^{(2)}(\mathbf{x})$ as the *activations* for a fixed value of \mathbf{x} . It follows that $h_k(\mathbf{x}, \mathbf{w}) = \sigma(a_k^{(2)}(\mathbf{x}))$. Taking the gradient of J with respect to the different sets of parameters $\mathbf{w}^{(1)}$ and $\mathbf{w}^{(2)}$, we find:

$$\nabla_{\mathbf{w}_{kj}^{(2)}} J(\mathbf{w}) = \frac{1}{h(\mathbf{x}, \mathbf{w})} \sum_k \sigma(a_k)(1 - \sigma(a_k)) \quad (6)$$

$$\nabla_{\mathbf{w}^{(2)}} J(\mathbf{w}) = \quad (7)$$

2.1.2 Stochastic Gradient Descent

2.2 Computational Results

2.2.1 Toy Problem

2.2.2 MNIST Data