

# Machine Learning 6.867 - Pset 3

November 6, 2015

## 1 Multi-Class SVM

## 2 Neural Networks

Neural networks are used in machine learning to make predictions, similar to logistic regression, SVM, or regression. We can represent neural networks using a graph with nodes and edges (see Bishop figure 5.1). Assume that we observe data  $(\mathbf{x}^{(i)}, \mathbf{y}^{(i)}), i = 1, \dots, N$ , where  $\mathbf{x}^{(i)} \in \mathbb{R}^D$  and  $\mathbf{y}^{(i)} \in \{0, 1\}^K$ . Let  $(\mathbf{x}, \mathbf{y}) = ([x_0, x_1, \dots, x_D], [y_1, \dots, y_K])$  be a general observation. We create nodes for each of the features  $x_i$ , referred to as *inputs*, and nodes for each of the class labels  $y_i$ , referred to as *outputs*. Next, we introduce a series of nodes in the middle of the graph, called *hidden units*, and we draw edges connecting the **inputs**  $\rightarrow$  **hidden units**  $\rightarrow$  **outputs**. The key idea in neural networks is that we can model the hidden units as functions of the inputs, and model the outputs as functions of the hidden units.

For 2-layer neural networks, we have one layer of hidden units denoted by  $[z_0, z_1, \dots, z_m]$ . There are weights  $\mathbf{w}_{ji}^{(1)}$  for each edge  $x_i \rightarrow z_j$  and  $\mathbf{w}_{kj}^{(2)}$  for each edge  $z_j \rightarrow y_k$ , which are unknown and will be learned through training the model. Let  $\sigma(\cdot)$  denote the logistic function, which we will use as the *activation function* for both the hidden and output layers of our neural network. The predicted value for output  $k$  given parameters  $\mathbf{w} = \{\mathbf{w}^{(1)}, \mathbf{w}^{(2)}\}$  and input  $\mathbf{x}$  will be:

$$h_k(\mathbf{x}, \mathbf{w}) = \sigma \left( \sum_{j=0}^M w_{kj}^{(2)} \sigma \left( \sum_{i=0}^D w_{ji}^{(1)} x_i \right) \right). \quad (1)$$

### 2.1 Implementation

#### 2.1.1 Gradient Descent

We implemented a 2-layer regularized neural network using gradient descent. The loss function, which is the negative log-likelihood, is equal to:

$$l(\mathbf{w}) := \sum_{i=1}^N \sum_{k=1}^K \left[ -y_k^{(i)} \log(h_k(\mathbf{x}^{(i)}, \mathbf{w})) - (1 - y_k^{(i)}) \log(1 - h_k(\mathbf{x}^{(i)}, \mathbf{w})) \right] \quad (2)$$

We add a regularizer term, and the final cost function becomes:

$$J(\mathbf{w}) := l(\mathbf{w}) + \lambda(\|\mathbf{w}^{(1)}\|_F^2 + \|\mathbf{w}^{(2)}\|_F^2), \quad (3)$$

where  $\|A\|_F = \sqrt{\sum_{i,j} A_{ij}^2}$  is the matrix Frobenius norm. Taking the gradient of  $J$  with respect to the different sets of parameters  $\mathbf{w}^{(1)}$  and  $\mathbf{w}^{(2)}$ , we find:

$$\nabla_{\mathbf{w}^{(1)}} J(\mathbf{w}) = \frac{1}{h(\mathbf{x}, \mathbf{w})} \sum_k \sigma(a_2)(1 - \sigma(a_2)) \quad (4)$$

$$\nabla_{\mathbf{w}^{(2)}} J(\mathbf{w}) = \quad (5)$$

### 2.1.2 Stochastic Gradient Descent

## 2.2 Computational Results

### 2.2.1 Toy Problem

### 2.2.2 MNIST Data